

Parallel Computing

Ekkapot Charoenwanit

Software Systems Engineering

TGGS

KMUTNB

Lecture 7:

□ Shared-Memory Programming with OpenMP

- Functional Parallelism
- Synchronization Constructs
- Thread-Local Storage

OpenMP: Sections Construct

The **sections** construct provides **functional parallelism** by enabling **different threads** to carry out **different tasks**:

- ❑ It permits us to specify several different code regions, each of which is represented by a **section** construct and is executed by one thread in the team.
- ❑ It consists of **two different directives** as follows:

```
#pragma omp sections [clause[,] clause]...  
{  
    [#pragma omp section ]  
        structured block  
    [#pragma omp section  
        structured block ]  
    ...  
}
```

OpenMP: Sections Construct

- ❑ The code block of each section must be a **structured block**, i.e. one entry and one exit.
- ❑ At runtime, one thread executes one section at a time, and one section will be executed **exactly once**.
- ❑ If there are **fewer** threads than sections, some or all threads must execute multiple sections.
- ❑ If there are **more** threads than sections, the remaining threads will be idle.

OpenMP: Sections Construct

```
#pragma omp parallel
{
    #pragma omp single
    std::cout << "The number of threads : " << omp_get_num_threads() << std::endl;

    #pragma omp sections
    {
        #pragma omp section
        f();
        #pragma omp section
        g();
        #pragma omp section
        h();
    }
}
```

If three or more threads are available, one thread might invoke *f()*, a second one might invoke *g()* and a third one might invoke *h()* whereas the remaining are idle, waiting at the implicit barrier.

OpenMP: Synchronization

Recall that we need to:

- ❑ synchronize actions on **shared variables**
- ❑ ensure **the correct ordering** of **reads** and **write** to shared variables
- ❑ protect **concurrent updates** to shared variables
 - Updates to variables are **not atomic by default**.

OpenMP: Barrier Construct

The OpenMP directive used to implement an **explicit barrier** in OpenMP is

```
#pragma omp barrier
```

- ❑ No threads cannot progress past a **barrier construct** until all threads in the team have arrived.
 - The all-or-none Principle !!!
- ❑ Either all threads in the same team or none must encounter a barrier.
 - Otherwise, this will result in a **deadlock**.
- ❑ This is one of the very few OpenMP directives that do not have an associated block of code.

OpenMP: Critical Section

In OpenMP, a **critical section** can be denoted by the **critical section construct** which provides a means to ensure that multiple threads do not attempt to update shared data simultaneously.

<pre>#pragma omp critical [(name)] structured block</pre>

- ❑ An **optional name** can be given to a **critical construct** and this name is **global**.
- ❑ When a thread encounters a critical section, it waits until no thread is executing inside the critical section with **the same name**.

OpenMP: Critical Section

```
#pragma omp parallel default(none) shared(totalHits) private(seedVal)
{
    seedVal = omp_get_thread_num();//Each thread seeds with its own thread id.

    #pragma omp for
    for(int i=0 ; i<NUMITERATIONS ; i++)
    {
        double x = (double) rand_r(&seedVal) / (double) RAND_MAX;
        double y = (double) rand_r(&seedVal) / (double) RAND_MAX;

        double result = std::sqrt((x*x) + (y*y));

        if(result<1.0){
            #pragma omp critical(update_total_hits)
            {
                totalHits += 1.0;//check if the generated value is inside a unit circle.
            }
        }
    }
}
```

In this snippet, the program calculates the value of π using Monte Carlo simulation.

- Updates to the shared variable *totalHits* are protected by the critical pragma *update_total_hits*.
- However, speedup is poor with this approach.
- It is more efficient to use the reduction clause.

OpenMP: Atomic Construct

The **atomic** directive is used to **protect** a **single update** to a **shared variable**:

- ❑ It applies to only a **single statement**, instead of a code block.
- ❑ The syntax of the **atomic** directive is

<pre>#pragma omp atomic <i>statement</i></pre>
--

where *statement* is a single arithmetic statement.

OpenMP: Single Construct

The **single** construct is a **work-sharing construct** that specifies that the associated code block should be executed by **one thread only**.

- ❑ There is always an **implicit barrier** at the end.
- ❑ The other threads wait at **the barrier** until the code block has been executed by the chosen thread.
- ❑ The **implicit barrier** at the end of the construct can be removed with the *nowait* clause.
- ❑ The syntax of the **single** directive is

<code>#pragma omp single [clause[[,] clause] ...]</code>	<code>structured block</code>
---	-------------------------------

OpenMP: Single Construct

- ❑ It is up to the OpenMP runtime which thread in the team to execute the code associated with a **single** construct.
- ❑ This construct can be useful when:
 - if one thread in the team is **an I/O thread**.
 - there is **a setup part** such as **memory allocation** and **deallocation** where only one thread must be responsible for.

OpenMP: Single Construct

```
private(list)  
firstprivate(list)  
copyprivate(list)  
nowait
```

Clauses supported by the **single** construct

OpenMP: CopyPrivate

The *copyprivate* clause provides a mechanism for **broadcasting** the value of a **private variable** from one thread to the other threads in the team:

- ❑ The typical use for this clause is to have one thread initialize or read private data that is subsequently used by the other threads as well.
- ❑ After the **single** construct has ended, but before the threads have left the associated **barrier**, the values of the private variables specified by the *copyprivate* clause are **copied** to the other threads.
- ❑ The standard **prohibits** the use of this clause in combination with the *nowait* clause.

OpenMP: CopyPrivate

```
int x = 1;
int y = 2;

#pragma omp parallel default(none) private(x) shared(y, std::cout)
{
    #pragma omp single copyprivate(x)
    {
        x = 10;
        y = y + 2*x;
    }

    #pragma omp critical
    std::cout << "Thread " << omp_get_thread_num() << " has x = " << x << " and y = " << y << std::endl;
}
```

OpenMP: Master Construct

The **master** construct is a synchronization construct.

- ❑ The syntax of the **master** directive.

```
#pragma omp master  
    structured block
```

- ❑ Only the master thread in the team executes the associated code block.
- ❑ It differs from the single construct in that there is no implicit barrier at the end.

OpenMP: Master Construct

```
#pragma omp parallel
{
    #pragma omp master
    {
        std::cout << "Parallel Region #1 : I'm Thread " << omp_get_thread_num() << std::endl;
    }
}

#pragma omp parallel
{
    if(omp_get_thread_num()==0)
        std::cout << "Parallel Region #2 : I'm Thread " << omp_get_thread_num() << std::endl;
}
```

The **master** construct can be implemented with an if-statement as shown in the code snippet above.

OpenMP: ThreadPrivate Clause

By default, global data is shared.

- ❑ However, in some situations, we may need or would prefer to have private data that persists throughout the computation.
- ❑ This is where the *threadprivate* clause comes in handy.
- ❑ The effect of the *threadprivate* clause is that the named global-lifetime objects are replicated so that each thread has its own copy.

OpenMP: ThreadPrivate Clause

❑ That is, each threads gets a **private** or a local copy of the specified **global variables**.

❑ The syntax is

```
#pragma omp threadprivate (list)
```

❑ Restriction is that

- The dynamic threads mechanism must be **turned off** with `omp_set_dynamic(false)` so that the number of threads across different parallel regions remains constant.

OpenMP: Nested Parallelism

If a thread in a team executing a parallel region encounters **another parallel region**, it creates a new team and becomes **the master thread** of that new team.

□ This is generally referred to in OpenMP as “**nested parallelism**”.

By **default**, this feature is **disabled**, meaning that the thread that encounters a nested parallel region will not form a new team, i.e., the encountered parallel region is **inactive**.

□ The `omp_get_nested()` routine can be used to test if nested parallelism is enabled.

OpenMP: Nested Parallelism

```
if(omp_get_nested())
    std::cout << "Nested Parallelism is ON." << std::endl;
else
    std::cout << "Nested Parallelism is OFF." << std::endl;

#pragma omp parallel
{
    #pragma omp critical
    std::cout << "Thread " << omp_get_thread_num() << " executes the outer parallel region." << std::endl;

    #pragma omp parallel
    {
        #pragma omp critical
        std::cout << "\tThread " << omp_get_thread_num() << " executes the inner parallel region." << std::endl;
    }/*--- End of Inner Parallel Region ---*/
}/*--- End of Outer Parallel Region ---*/
```

OpenMP: Nested Parallelism

Executing the code snippet on the previous slide might give you output that looks similar to the following output:

```
Nested Parallelism is OFF.  
Thread 3 executes the outer parallel region.  
Thread 0 executes the outer parallel region.  
    Thread 0 executes the inner parallel region.  
    Thread 0 executes the inner parallel region.  
Thread 1 executes the outer parallel region.  
    Thread 0 executes the inner parallel region.  
Thread 2 executes the outer parallel region.  
    Thread 0 executes the inner parallel region.
```

To enable **nested parallelism**, the `OMP_NESTED` environment variable needs to be set to **true**:

- ❑ This `OMP_NESTED` environment variable is currently **deprecated** by the standard.
- ❑ In a future version of OpenMP, you might need to use a different mechanism for enabling and disabling nested parallelism.

OpenMP: Nested Parallelism

```
Nested Parallelism is ON.  
Thread 3 executes the outer parallel region.  
Thread 0 executes the outer parallel region.  
Thread 2 executes the outer parallel region.  
    Thread 0 executes the inner parallel region.  
    Thread 2 executes the inner parallel region.  
    Thread 1 executes the inner parallel region.  
    Thread 3 executes the inner parallel region.  
Thread 1 executes the outer parallel region.  
    Thread 0 executes the inner parallel region.  
    Thread 1 executes the inner parallel region.  
    Thread 2 executes the inner parallel region.  
    Thread 3 executes the inner parallel region.  
    Thread 3 executes the inner parallel region.  
    Thread 1 executes the inner parallel region.  
    Thread 0 executes the inner parallel region.  
    Thread 2 executes the inner parallel region.  
    Thread 2 executes the inner parallel region.  
    Thread 1 executes the inner parallel region.  
    Thread 3 executes the inner parallel region.  
    Thread 0 executes the inner parallel region.
```

After enabling **nested parallelism** by setting the *OMP_NESTED* environment variable to **true**.

OpenMP: Collapse Clause

A **short loop** (one with a small number of iterations) has a downside in a serial program:

- ❑ The **loop distribution overhead** might dominate the performance if there is **no sufficient work** to do per loop iteration.
- ❑ A short loop might not only limit the number of threads that may be used, but also leads to **load imbalance** if the number of iterations is not a multiple of the number of threads.

For **perfectly nested rectangular loops**, we can parallelize **multiple loops** in the **loop nest** using the ***collapse*** clause by **merging** or “**collapsing**” the multiple loops in the loop nest into one single and longer loop.

References

- [1] *Barbara Chapman, Gabriele Jost, and Ruud van der Pas. 2007. Using OpenMP: Portable Shared Memory Parallel Programming (Scientific and Engineering Computation). The MIT Press.*
- [2] *Robit Chandra, Leonardo Dagum, Dave Kohr, Dror Maydan, Jeff McDonald, and Ramesh Menon. 2001. Parallel programming in OpenMP. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA.*