

Parallel Computing

Ekkapot Charoenwanit

Software Systems Engineering

TGGS

KMUTNB

Lecture 8:

□ Shared-Memory Programming with OpenMP

➤ OpenMP Tasks

OpenMP: Tasks

OpenMP tasks provide a new parallel programming paradigm

- ❑ called the work-oriented paradigm
- ❑ based on the concept of a task pool

In this work-oriented paradigm, units of work (referred to as **OpenMP tasks**)

- ❑ are generated or “**pushed**” into the task pool by threads
- ❑ are retrieved or “**pulled**” off the task pool and executed by threads

OpenMP: Tasks

In **classic OpenMP**, threads are treated as a fundamental concept:

- ❑ The thread-centric paradigm
- ❑ The OpenMP standard *prior to 3.0* described semantics in terms of threads.

In the work-oriented paradigm, we focus on **units of work** referred to as **tasks**:

- ❑ We must now think how the code and data can be broken into units of work that can be executed in parallel, i.e., **tasks**.
 - ❑ We can think of a task as
 - code
 - data (Data Environment)
 - internal control variables (ICVs) [**implementation details**]
- packaged up as an independent schedulable unit.

OpenMP: Tasks

Threads are assigned to perform the work of each task:

- ❑ Tasks may be **deferred**.
- ❑ Tasks may be **executed** immediately.

Tasks enable **irregular** and **recursive** computational problems to be parallelized in a natural way, e.g.,

- ❑ Traverse a linked list while performing work on each node in parallel
- ❑ Implement parallel recursive algorithms

OpenMP: Tasks and Threads

- ❑ A task is a specific instance of **Executable Code** and its **Data Environment** that can be scheduled for execution by a thread.
- ❑ Each encountering thread creates a new instance of a task, e.g. **code** and **data**.
 - Tasks can be **deferred**, i.e., they do not need to be **executed** immediately.
 - If a task is deferred, some thread in the team executes the task at some time later.
 - The encountering thread that generated a task is not necessarily the same thread that will eventually execute the task.
- ❑ Tasks bind to the corresponding **innermost** enclosing parallel regions.

OpenMP: Task Creation


Tasks are created when

- ❑ when a `#pragma omp parallel` is reached
 - implicit tasks are created per thread
- ❑ when a `#pragma omp task` is encountered
 - an explicit task is created
- ❑ when a `#pragma omp taskloop` is encountered
 - explicit tasks per chunk are created
- ❑ when a `#pragma omp target` is encountered
 - A target task is created

C/C++

```
#pragma omp task [clause]  
... structured block ...
```

OpenMP: Task Creation and Execution



OpenMP: Task Creation and Execution

```
#pragma omp parallel num_threads(2)
{
    #pragma omp task //Task A
    {
        #pragma omp critical
        std::cout << "Thread " << omp_get_thread_num() << " executes Task A." << std::endl;
    }

    #pragma omp task //Task B
    {
        #pragma omp critical
        std::cout << "Thread " << omp_get_thread_num() << " executes Task B." << std::endl;
    }
}

/*--- End of Parallel Region ---*/
```

```
Thread 0 executes Task A.
Thread 0 executes Task B.
Thread 0 executes Task A.
Thread 1 executes Task B.
```




- Each thread in `#pragma omp parallel` creates a new task.
- It is non-deterministic which threads are to execute which tasks.

OpenMP: Task Creation and Execution

```
#pragma omp parallel num_threads(2)
{
    #pragma omp task //Task A
    {
        #pragma omp critical
        std::cout << "Thread " << omp_get_thread_num() << " executes Task A." << std::endl;
    }

    #pragma omp task //Task B
    {
        #pragma omp critical
        std::cout << "Thread " << omp_get_thread_num() << " executes Task B." << std::endl;
    }
}
/*--- End of Parallel Region ---*/
```



- ❑ Each thread inside `#pragma omp parallel` pushes two tasks to the task pool, i.e., four tasks are pushed to the task pool in total.
- ❑ There is an implicit barrier at the end of `#pragma omp parallel`, which acts as a **task synchronization construct**.

OpenMP: Task Creation and Execution

```
#pragma omp parallel num_threads(2)
{
    #pragma omp single
    {
        #pragma omp task //Task A
        {
            #pragma omp critical
            std::cout << "\tThread " << omp_get_thread_num() << " executes Task A." << std::endl;
        }

        #pragma omp task //Task B
        {
            #pragma omp critical
            std::cout << "\tThread " << omp_get_thread_num() << " executes Task B." << std::endl;
        }
    }
}/*--- End of Single Construct ---*/
}/*--- End of Parallel Region ---*/
```

Thread 1 executes Task A.
Thread 0 executes Task B.



- One of the two threads generates Task A and Task B.
- It is non-deterministic which threads are to execute which tasks.

OpenMP: Task Creation and Execution

```
#pragma omp parallel num_threads(2)
{
    #pragma omp single
    {
        #pragma omp task //Task A
        {
            #pragma omp critical
            std::cout << "\tThread " << omp_get_thread_num() << " executes Task A." << std::endl;
        }

        #pragma omp task //Task B
        {
            #pragma omp critical
            std::cout << "\tThread " << omp_get_thread_num() << " executes Task B." << std::endl;
        }
    }
}/*--- End of Single Construct ---*/
}/*--- End of Parallel Region ---*/
```



- ❑ One of the two threads inside `#pragma omp parallel` is in charge of enqueueing tasks to the task pool.
- ❑ There is an implicit barrier at the end of `#pragma omp single`, which acts as a task synchronization construct.
- ❑ The *nowait* clause can be used on `#pragma omp single` to remove the implicit barrier.

OpenMP: Task Execution Model

Task Creation:

OpenMP tasks can be created using the following **task creation patterns**:

- ☐ single task creator
 - with **#pragma omp single**
- ☐ multiple task creators
- ☐ nested tasks

Task Execution:

All threads in the team are candidates to execute tasks:

- ☐ **deferred**
- ☐ **Immediately executed**

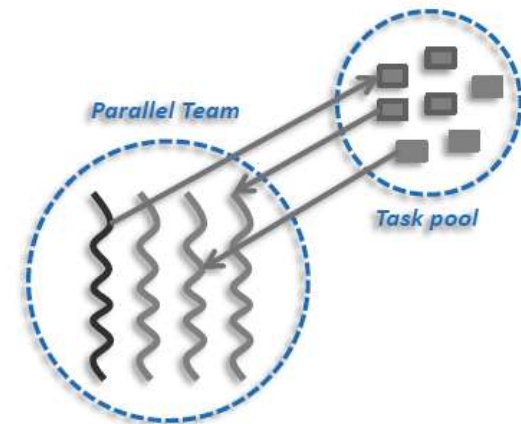


Figure Credit: C. Terboven PPCES 2021

OpenMP: Task Clauses

→ private(list) → firstprivate(list) → shared(list) → default(shared none) → in_reduction(r-id: list)	Data Environment
→ allocate([allocator:] list) → detach(event-handler)	Miscellaneous

→ if(scalar-expression) → mergeable → final(scalar-expression)	Cutoff Strategies
→ depend(dep-type: list)	Synchronization
→ untied → priority(priority-value) → affinity(list)	Task Scheduling

Figure Credit: C. Terboven PPCES 2021

OpenMP: Data Environment

- ❑ The Data Environment consists of all the variables associated with the execution of a given task.
 - Because tasks can be *deferred*, the data is “*captured*” at task creation.
- ❑ The semantics of data scopes are adapted to deferred execution as follows:
 - For a *shared* variable, the reference to the *shared* variable within the task is to the memory location to that name at the time where the task was created.
 - For a *private* variable, the reference to the *private* variable within the task is to new *uninitialized storage* that is created when the task is executed.
 - For a *firstprivate* variable, the reference to the *firstprivate* is to new storage that is created and initialized with the value of the existing memory of that name when the task is created.
 - Capturing data structures with *firstprivate* captures only *the pointer* but *not the pointed data*.

OpenMP: Data Environment

Data-Scoping Rules: Most rules from `#pragma omp parallel` still apply.

Pre-determined data-sharing rules for `#pragma omp task`:

- ☐ Static and global variables are *shared*.
- ☐ Static local variables are *shared*.
- ☐ Automatic (local) variables are *private*.
- ☐ *Threadprivate* variables are *threadprivate*.

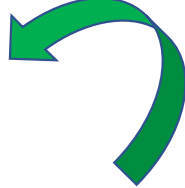
The data-sharing attributes of variables that are not listed in any data-sharing clauses of `#pragma omp task` and are *not pre-determined* according to the rules above, are implicitly determined as follows:

Implicit data-sharing rules for `#pragma omp task`:

- ☐ The *shared* attribute is lexically inherited.
- ☐ In any other case, the variable is *firstprivate*.

OpenMP: Data Environment

```
int a = 1;
void foo()
{
    int b = 2, c=3;
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;
            //a is shared:      a = 1
            //b is firstprivate: b = undefined
            //c is shared:      c = 3
            //d is firstprivate: d = 4
            //e is private:     e = 5
        }
    }
    /*--- End of Parallel Region ---*/
}
```



- ❑ In `#pragma omp parallel`, `c` is shared by default
 - In `#pragma omp task`, the data scoping of `c` is lexically inherited, i.e., `c` is *shared*.
- ❑ In `#pragma omp parallel`, `a` is shared by default since it is a global variable.
 - In `#pragma omp task`, the data scoping of `a` is lexically inherited, i.e., `a` is *shared*.

OpenMP: Task Scheduling

Tasks can be either *tied* or *untied*:

- ❑ Tasks are *tied* by default unless the *untied* clause is explicitly placed on `#pragma omp task`.
- ❑ Tasks remain *tied* to the threads that first execute them.

A *tied* task is *always executed* by *the same thread*:

- ❑ The *tied* task must be executed by the same thread throughout its entire life cycle.
- ❑ The thread that executes the *tied* task does not have to be the same thread that created the task.

Tied tasks can restrict scheduling and may, therefore, negatively impact performance.

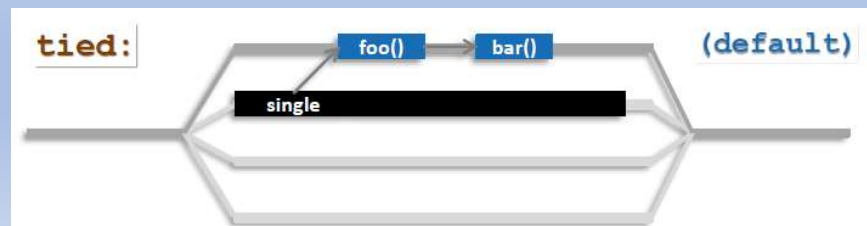


Figure Credit: C. Terboven PPCES 2021

OpenMP: Task Scheduling

In addition to **explicit tasks** specified with **#pragma omp task**, OpenMP has also the notion of **implicit tasks**.

An implicit task is one that is **implicitly created** when **a parallel construct** specified with **#pragma omp parallel** is encountered.

- ❑ Each implicit task is assigned to a different thread in the team and remains **tied** to that thread throughout the entire life cycle of the task.

OpenMP: Task Scheduling

The scheduling constraint can be relaxed using the *untied* clause:

- ❑ An *untied* task can migrate between any threads in the team.
- ❑ An *untied* task must be avoided when used with *thread-centric features* such as
 - Thread ID
 - Threadprivate Data [[See](#)]
- ❑ The OpenMP runtime has *greater flexibility* in scheduling tasks.
 - can potentially provide *better load-balancing*

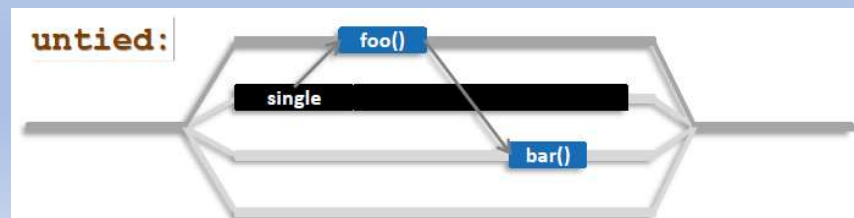


Figure Credit: C. Terboven PPES 2021

OpenMP: Task Scheduling

Task Scheduling Points (TSP):

- ❑ Tasks can be **suspended** and **resumed** at TSPs.
- ❑ **Implicit TSPs** are included at the following locations:
 - Task creation
 - the point of encountering **#pragma omp task**
 - Task synchronization
 - The point of encountering **#pragma omp taskwait**
 - The end of **#pragma omp taskgroup**
 - The completion point of a task
 - The point of encountering an implicit or an explicit barrier
- ❑ **Explicit TSPs** can be specified using **#pragma omp taskyield**.

C/C++

```
#pragma omp taskyield
```

OpenMP: Task Scheduling

At a **TSP**, a thread that is executing a task *A* may

- ❑ temporarily suspend **the current task** *A* and
- ❑ switch to execute **a different task** *B*
 - It can either start or resume this other task *B*.

OpenMP: Task Scheduling

The taskyield directive explicitly specifies that **the current task** can be **suspended** in favor of **a different task**.

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task untied
    {
        foo();
        #pragma omp taskyield
        bar()
    }
}
```




Figure Credit: C. Terboven PPCES 2021

OpenMP: Task Scheduling

Caveat with threadprivate and thread-specific information:

- ❑ When a thread encounters a TSP, the OpenMP implementation may choose to **suspend the current task** and **schedule the thread to work on another task**.
 - This implies that the value of a **threadprivate** variable or other thread-specific information, e.g. **Thread ID**, can potentially change across a TSP.
- ❑ If the suspended task is a **tied** task, then the thread that resumes executing the task is guaranteed to be the same thread that suspended it.
 - The **Thread ID** will remain the same after the task is resumed.
- ❑ If the suspended task is an **untied** task, the thread that resumes executing the task may be different from the thread that suspended it.
 - Both the **Thread ID** and the value of a **threadprivate** variable before and after the TSP may be different.

OpenMP: Task Scheduling

Let's sum up the concept of **tiedness** and **untiedness**.

❑ When a thread encounters a TSP, it may do one of the following:

- it may begin execution of a **tied** task bound to the current team
- it may resume execution of any suspended **tied** task, bound to the current team, to which it is **tied**.
- it may begin execution of an **untied** task bound to the current team
- it may resume execution of any suspended **untied** task bound to the current team

❑ If more than one of the above choices is available, it is **unspecified** as to which option will be chosen by the OpenMP runtime.

OpenMP: Task Synchronization

- ❑ All tasks created by any thread of the current team are guaranteed to have completed at a barrier (implicit or explicit).
 - ❑ Explicit barrier : `#pragma omp barrier`
 - ❑ Implicit barrier : at the end of
 - `#pragma omp parallel`
 - `#pragma omp for`
 - `#pragma omp single`
 - etc.
- ❑ A task that encounters `#pragma omp taskwait` is suspended until **all child tasks** complete.
 - applies only to **child tasks, not all descendant tasks !!!**
 - provides **shallow task synchronization**.

C/C++

```
#pragma omp taskwait
```

OpenMP: Task Synchronization

```
1  int main(int argc,  
2      char* argv[])  
3  {  
4      [...]  
5      #pragma omp parallel  
6      {  
7          #pragma omp single  
8          {  
9              fib(input);  
10         }  
11     }  
12     [...]  
13 }
```

```
14 int fib(int n) {  
15     if (n < 2) return n;  
16     int x, y;  
17     #pragma omp task shared(x)  
18     {  
19         x = fib(n - 1);  
20     }  
21     #pragma omp task shared(y)  
22     {  
23         y = fib(n - 2);  
24     }  
25     #pragma omp taskwait  
26     return x+y;  
27 }
```

- ❑ Only one thread executing `#pragma omp single` is responsible for creating tasks.
- ❑ This thread creates the two initial tasks (line 17 and 21).
- ❑ `#pragma omp taskwait` (line 25) is used to wait for the two tasks to complete. Otherwise, x and y would get lost.

OpenMP: Task Synchronization

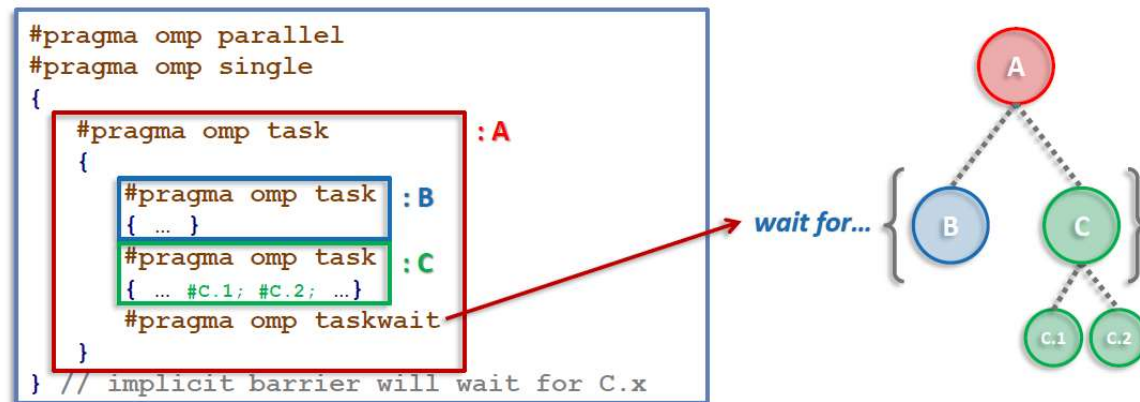


Figure Credit: C. Terboven PPCES 2021

- ❑ Task *C.1* and *C.2* are not affected:
 - They can execute past `#pragma omp taskwait` without the need to wait.
 - `#pragma omp taskwait` only waits for Task *B* and *C*.
- ❑ All tasks including *C.1* and *C.2* are guaranteed to have completed at the implicit barrier of the `#pragma omp single`.

OpenMP: Task Synchronization

Deep Task Synchronization:

- ❑ Tasks can be synchronized across all descendant levels can be synchronized with `#pragma omp taskgroup`.
 - provides **deep task synchronization**.
- ❑ The clause allowed is the `task_reduction` clause:
 - of the form `task_reduction(reduction-identifier:list-items)`

```
#pragma omp taskgroup [clause[,] clause]...  
{structured-block}
```

OpenMP: Task Synchronization

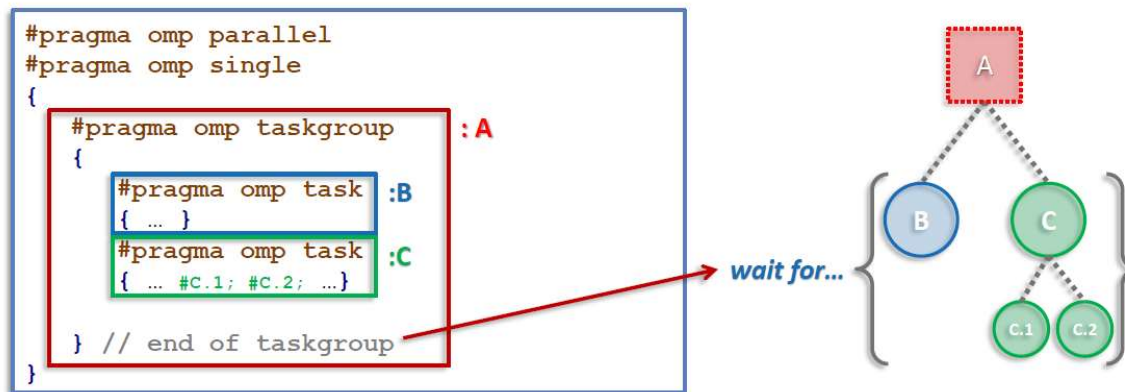


Figure Credit: C. Terboven PPCES 2021

- ❑ Task *C.1* and *C.2* are affected:
 - They **can not** execute past `#pragma omp taskwait` without the need to wait.
- ❑ All tasks, namely, *B*, *C*, *C.1* and *C.2* are guaranteed to have completed at the end of the `#pragma omp taskgroup`.

OpenMP: Task Reduction

```
int res = 0;
node_t* node = NULL;
...
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp taskgroup task_reduction(+: res)
        { // [1]
            while (node) {
                #pragma omp task in_reduction(+: res) \
                    firstprivate(node)
                { // [2]
                    res += node->value;
                }
                node = node->next;
            }
        } // [3]
    }
}
```

- ❑ The taskgroup-scoping reduction *task_reduction* clause
 - registers a new reduction at [1].
 - computes and make the final result available after [3].
- ❑ The task *in_reduction* clause
 - allows the task to participate in the reduction operation [2]
- ❑ introduced in OpenMP 5.0

Figure Credit: C. Terboven PPCES 2021

OpenMP: Cut-Off Strategies

For recursively defined divide-and-conquer algorithms, a *cut-off* can be defined to avoid *task explosion*.

Mechanisms:

- ❑ Programmed cut-off mechanism by explicitly calling a function that does not create tasks
- ❑ The *if* clause
- ❑ The *final* clause
- ❑ The *mergeable* clause

OpenMP: Programmed Cut-Off

```
float sum(const float* a, size_t n)
{
    // base cases
    if(n==0) return 0;
    if(n==1) return *a;

    // recursive case
    size_t half=n/2;
    return sum(a, half) + sum(a+half, n-half);
}
```

OpenMP: Programmed Cut-Off

```
float sum(const float* a, size_t n)
{
    // base cases
    if(n == 0) return 0;
    if(n == 1) return *a;

    // recursive case
    size_t half = n/2;
    float x, y;
    #pragma omp parallel
    #pragma omp single nowait
    {
        #pragma omp task shared(x)
        x = sum(a, half);
        #pragma omp task shared(y)
        y = sum(a+half, n-half);
        #pragma omp taskwait
        x += y;
    }
    return x;
}
```

- ❑ The OpenMP code creates tasks recursively.
- ❑ Problem:
 - The code creates too many small tasks, i.e., *n* is small.
 - Inefficient
- ❑ Solution:
 - stop creating tasks at some *cut-off threshold*

OpenMP: Programmed Cut-Off

```
#define CUTOFF 100 // arbitrary

static float parallel_sum(const float *a,
    size_t n){
    // base case
    if (n <= CUTOFF){
        return serial_sum(a, n);
    }
    // recursive case
    float x, y;
    size_t half = n / 2;
    #pragma omp task shared(x)
    x = parallel_sum(a, half);
    #pragma omp task shared(y)
    y = parallel_sum(a + half, n - half);
    #pragma omp taskwait
    x += y;
    return x;
}
```

```
float sum(const float *a, size_t n){
    float r;
    #pragma omp parallel
    #pragma omp single nowait
    r = parallel_sum(a, n);
    return r;
}

static float serial_sum(const float *a,
    size_t n){
    float x = 0.0;
    for(int i=0; i<n; i++) {
        x += a[i];
    }
    return x;
}
```

Drawback: separate versions of sum !!!

- ☐ serial_sum
- ☐ parallel_sum

OpenMP: If Clause

- ❑ If the *if* clause evaluates to *false*
 - the encountering task, i.e., *parent*, is suspended
 - the new task, i.e., *child*, is executed immediately
 - The new task is denoted as an “*undelayed*” task
 - the parent task resumes once the new task has completed
- ❑ This feature can be used to improve the performance by avoiding queuing tasks that are *too small*.
 - The computational complexity of such small tasks may be considered too small and not worth the overhead for queuing them in the task pool.
- ❑ The *if* clause *does not affect* the child tasks created by the encountering task.
- ❑ This feature can be useful for debugging task-based OpenMP applications.

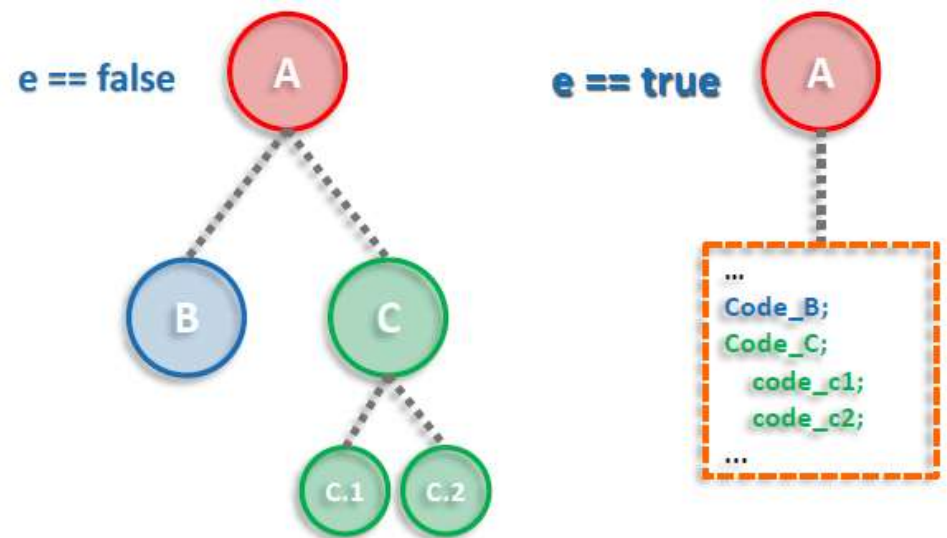
OpenMP: Final Clause

- ❑ The *final* clause prevents the creation of further tasks if condition evaluates to *true*.
- ❑ The purpose is to *reduce overhead*, which will *also limit the degree of parallelism* by preventing too many tasks from being created.
- ❑ The new task is created and executed normally
 - However, in its context, its nested tasks will be *executed immediately by the same threads*
 - The undeferred nested tasks are denoted as *"included"* tasks.

OpenMP: Final Clause

```
#pragma omp task final(e)
{
    #pragma omp task
    { ... }
    #pragma omp task
    { ... #C.1; #C.2 ... }
    #pragma omp taskwait
}
```

Figure Credit: C. Terboven PPCES 2021



OpenMP: Mergeable Clause

- ❑ A task can be created mergeable if
 - it is an *undelayed* task **OR**
 - it is an *included* task
- ❑ Its purpose is to
 - reduce OpenMP runtime overhead
 - inline the Data Environment to enable efficient data sharing
- ❑ Often used with the *final* clause
- ❑ No commercial OpenMP implementations have taken advantage of *mergeable* yet.

Reference

- [1] *Ruud van der Pas, Eric Stotzer, and Christian Terboven. 2017. Using OpenMP -- The Next Step: Affinity, Accelerators, Tasking, and SIMD (1st. ed.). The MIT Press.*
- [2] *Christian Terboven. 2021. Programming OpenMP. PPCES 2021.*